

The Feature Ledger

● VERSION 2 · DRAFT, NOT YET PRESENTED

Everything feature-ledger does today, in plain language. Each entry is one capability you could ask for by name, sized so the biggest pieces stand out.

What's changed since the last edition (September 8, 2026): every entry below is highlighted to show what's different. A green stripe marks a brand-new feature; a blue stripe marks an existing feature that was materially reworked, with its own **What changed** note. Anything without a stripe hasn't changed.


New feature

Updated feature

THE LEDGER AT A GLANCE

23 features across 4 areas

EVERY ONE OF THOSE 23, BY SIZE



SINCE THE LAST EDITION, SEPTEMBER 8, 2026

1 NEW 3 UPDATED 19 UNCHANGED

CONTENTS

01	Setting up a ledger	6	03	Reading the ledger	3
02	Recording what the product does	9	04	The documents it produces	5

01 Setting up a ledger

6 FEATURES · 2 UPDATED

BIG

UPDATED

Creating a ledger in an existing project

previously “One-command setup in any codebase”

The command `ledger init` creates a `.ledger` directory holding the whole corpus: the product name, the accent colour and the style guide you chose, and version 1 open as the release in progress. Everything the tool needs sits in that one directory, so a Python, Go or Rails repo looks identical to it. Setup also settles where the generated documents go. The Markdown docs are build output and are added to `.gitignore`. The PDF editions for the client are build output too, unless you run `ledger init --commit-pdfs`, which keeps them in `docs/client` so every document handed over stays in the repo.

WHAT CHANGED

- The client PDFs were written inside the gitignored output directory, with no way to keep them. A team that hands a printed edition to a client at each release wants the repo to hold what was handed over, so where the editions land is now a choice made once at setup.

BIG

The tone every entry is written in

A ledger is written in one tone, chosen from four style guides: Google developer documentation, GOV.UK, Federal Plain Language, and ASD-STE100 Simplified Technical English. The command `ledger style` prints that tone as a sample paragraph of product prose and three feature descriptions written in it, which is the thing a writer or an agent matches. Running `ledger style set` moves the project to another guide, or to a tone you write as prose yourself in `.ledger/STYLE.md`.

 MEDIUM

Wiring into a coding agent

Setup appends a short stanza to whichever agent instruction files the repo already has — CLAUDE.md, AGENTS.md, GEMINI.md, .cursor/rules, .github/copilot-instructions.md or .windsurfrules — and creates AGENTS.md when there are none. The stanza is a pointer rather than a copy: it names six commands, the tone the project writes in, and the rule that a feature change is not finished until the ledger records it. The --agents flag chooses the files, and changing the tone rewrites the stanza in every one of them.

 MEDIUM

UPDATED

Baseline survey of an existing codebase

The command ledger bootstrap prints a brief you paste into a coding agent, which then reads the codebase and records what it already does as version 1. The brief carries the entry bar, the project's tone, and five steps: list the capabilities, decide the areas, write the entries, cut the baseline from the date of the first commit, and hand back the entries it was least sure of. Because reading the whole codebase against the record is exactly what an audit does, the brief also records the survey as a full sweep, which is what gives the release gate something to measure against from the first day. The agent runs it start to finish without stopping to ask you anything.

WHAT CHANGED

- The survey ended at the baseline edition, so a freshly bootstrapped project reported its codebase as never checked and the release gate stayed quiet until somebody ran an audit by hand. The survey had in fact just swept everything, so the brief now records that before cutting — the one moment when what the ledger holds and what the codebase does are known to agree. Pass --no-audit for a survey that is deliberately partial.

 SMALL

Moving an existing ledger to a new tone

The command ledger style rewrite prints a brief for rewriting a ledger that is already written into the tone the project has moved to. The agent works one area at a time and replaces the words in place, so nothing lands on the release and no entry is marked as changed.

The authoring rules

The command `ledger rules` prints the full rules for writing the ledger: what each field means, where the bar for an entry sits, and how to record a change, a rename or a removal. The project's tone is printed inside it, and a `RULES.md` in the ledger directory replaces the packaged text with your own house conventions.

02 Recording what the product does

9 FEATURES · 1 NEW

Recording a capability

The command `ledger add` records one capability, with the audience it matters to, its size, its name and a present-tense description of what the product does. The id you give it is lower-case kebab-case and never changes, so it carries the feature's history through every later rename. An entry filed under an area or a sub-section that doesn't exist is refused as you write it, and the message lists the ones that do.

Recording a change to a feature

The command `ledger update` replaces a feature's description with its full current-state text, and takes a list of change bullets alongside it. Each bullet says what the feature was before and why it moved, which is the part a reader cannot get from the description. A feature that existed before this release cannot be updated without one, so an edition never carries a changed entry with no reason attached.

 BIG

NEW

The ledger audit

The command `ledger audit` reads the codebase for capabilities the ledger never recorded, which is what happens whenever someone ships without a coding agent that read the project's instructions. It prints a brief for a coding agent rather than guessing: the range since the last sweep, the changed paths already grouped into batches, and the commit subjects behind them. A sweep covers everything since the last audit, or the whole codebase with `--full`, which is the only way to find a capability the very first survey missed. The agent records what it finds and then stamps the commit it audited, so from that point the tool knows how much of the repository has never been examined. Cutting a release with unswept work behind it is refused, and says how many commits and files are outstanding. A project that started with ledger bootstrap is already stamped, because the baseline survey is itself a full sweep.

 MEDIUM

Rewording an entry

The command `ledger reword` replaces the words of an entry in place. Nothing is recorded against the release and the entry is not marked as changed, so the improved wording reads as though it had always been there. It reaches back to fix the wording of an edition already printed, and the same command exists for a sub-section's heading and overview.

 MEDIUM

Grouping features into areas

Areas group the entries, and the order of the list is the reading order of every document the ledger produces. The command `ledger categories` add places one, before or after any other; `rename` and `remove` move the structure without touching the entries underneath it. A ledger with no areas puts everything in one called `Features` and prints it without a heading, so a small product reads as a list of what it does rather than as a chapter of one.

 SMALL

Renaming a feature

The command `ledger rename` changes what a feature is called, and every document footnotes the name it went by before. A feature that existed before this release needs a reason: one bullet saying what it was called and why the name moved.

 SMALL

Retiring a feature

The command `ledger remove` marks a feature as gone from this release onward, with the reason you give it. It appears once, under the heading `No longer available as of this edition`, and every earlier edition still shows the feature as it stood at the time.

 SMALL

Recording a product-wide change

The command `ledger other-change` records something that runs across the whole product, such as a wording pass over many screens or a change in what a colour means. It prints once at the end of the edition, under the heading `Also since last time`, instead of forcing the same note onto a dozen unrelated entries.

 SMALL

Sub-sections inside an area

A sub-section is a heading between an area and a feature, for an area with a real part inside it. Each one carries a name and a one-or-two-sentence overview of that part, and you file a feature into it when you record or update the feature. When every feature in a sub-section arrives in the same release, the edition flags the sub-section once rather than printing a column of identical `New` marks.

03 Reading the ledger

3 FEATURES

 MEDIUM

Listing and reading entries

The command `ledger list` prints the whole corpus, one short line per feature: its size, its id, its name, and whether it is new or changed in the release in progress. It filters by area, audience, size, or by what has moved this cycle. The command `ledger show` prints the single entry you are about to edit, and prints every version it has been through when you ask for its history.

 MEDIUM

Checking the ledger is sound

The command `ledger status` answers three questions at once: what the open release has collected so far, what the corpus holds, and how the setup resolves — the derived palette, the masthead, the fonts, the tone in force, and the browser that will print. Anything that is valid but would still print badly comes out as a warning at the end, such as an entry with no size or an area with nothing in it. It exits with an error when the corpus is broken and stays silent when there is nothing to say, so a commit hook can run it on every commit.

 SMALL

Exporting the whole ledger

The command `ledger export` prints the entire corpus as a single JSON document: the release timeline, the areas, the sub-sections, the product-wide changes, and every feature with its full history. It writes to a file instead when you name one, for a backup, a migration, or handing the data to something that isn't this tool.

04 The documents it produces

5 FEATURES · 1 UPDATED

 BIG

Cutting a release

The command `ledger release cut` gives the release in progress its name and date and opens the next one behind it, which is where everything recorded from then on lands. The release just cut becomes an edition that never changes again, and `ledger release list` shows the timeline with the number of entries each release covers. Cutting a release with nothing recorded against it is refused, because the edition would be identical to the one before it.

 BIG

The printed edition for the client

The command `ledger build` prints a PDF of one release: a front page carrying the masthead, the feature count, a bar breaking those features down by size, and a note of what changed since the last edition, then every entry in reading order. The document is real HTML and CSS, printed through the Chrome or Edge already on the machine, with the two fonts embedded so text copies cleanly out of the finished file. Any edition can be reprinted at any time, and it reproduces exactly what was presented that day.

Three Markdown documents for the team

The command ledger build also writes three Markdown documents from the same corpus: FEATURES.md, the client-facing list; FEATURES_EXTENDED.md, the same list with implementation notes under each entry; and DEV_FEATURES.md, the entries only the team sees. All three describe the latest release, so they read as what the product does today. They are build output: the corpus is what you commit, and it carries the documentation's history on its own.

New and changed highlighting

previously "What moved since the last edition"

Every edition marks its own changes against the edition before it: a green stripe for a feature that is new, a blue stripe and a "What changed" note for one that was reworked, a grey stripe for one that was already part of the product and is only now listed, and no stripe for an entry that stands as it did. The front page counts how many of each. A first edition has nothing behind it to compare against, so it prints as a plain inventory with no stripes anywhere on it.

WHAT CHANGED

- There were two stripes, and an entry an audit turned up had to print as one of them. A capability that had been in the product for two years then read as this sprint's work, which is the one thing a client would take issue with. The third stripe says what is actually true: the product did not move this cycle, the record did. It is a neutral grey rather than a fourth hue because green and blue are already spoken for, an accent in that band is refused for the same reason, and a fifth hue would leave a brand only reds, oranges, purples and browns to choose from.

Branding the documents for a client

One file, .ledger/brand.json, is the swap surface for a new client: a logo, one accent colour, and the two font families. The other four tints of the accent are derived from the hex you give it, and a ledger with no logo sets the product name as a serif wordmark on the front page. The accent cannot be green, blue or teal, because those hues already mean new and reworked in every edition, and ledger status warns when a brand picks one that collides.